

A

1)

Can our customers be grouped into separate categories using k-means clustering so we may target them with different advertising and promotional campaigns?

2

One goal would be to define at least two groups of customers from the churn data set using k-means clustering.

B

1)

K-means clustering analyzes the data set by partitioning a data set into a specified number of clusters. Cluster centers are randomly chosen from points in the data set. All other data points are then assigned to the cluster that is closest to it. Cluster centers are recalculated using the means of the data points assigned to it. Data points are assigned to the clusters again. This process can repeat many times. An expected outcome is that all data points will be assigned to a cluster and the clusters will have minimum inertia when the algorithm has finished.

2)

One assumption of k-means clustering is that the created clusters are spherical.

3)

numpy

This is used for working with numpy arrays that are returned from the scaling function.

matplotlib.pyplot

This is used for visualizing inertia using the elbow method to find the optimal number of clusters.

sklearn.cluster import KMeans

This is the actual clustering algorithm that creates the model to cluster our data.

```
from sklearn.preprocessing import StandardScaler
```

This will be used to standardize the continuous variables.

C

1)

One data preprocessing goal is to standardize the variables before clustering them.

2) The initial data set variables will be:

```
Age continuous,  
Income continuous,  
Outage_sec_perweek continuous,  
MonthlyCharge continuous,  
Bandwidth_GB_year continuous,  
Contacts continuous,  
Yearly_equip_failure continuous,  
Tenure continuous
```

In []:

3)

prepare data

read in data and drop index column.

```
In [1]: import pandas as pd  
# Assuming your CSV file is named 'data.csv', adjust the file path as needed  
file_path = '/home/dj/skewl/D212/1/churn_clean.csv'  
pd.set_option('display.max_columns', None)  
# Read the data from the CSV file into a DataFrame  
df = pd.read_csv(file_path)  
#drop index column  
df = df.loc[:, ~df.columns.str.contains('Unnamed')]
```

check for missing values.

```
In [2]: # Identify missing values using isna() method  
missing_values = df.isna().sum()  
# Print DataFrame with True for missing values and False for non-missing values  
print(missing_values)  
# no missing values.
```

CaseOrder	0
Customer_id	0
Interaction	0
UID	0
City	0
State	0
County	0
Zip	0
Lat	0
Lng	0
Population	0
Area	0
TimeZone	0
Job	0
Children	0
Age	0
Income	0
Marital	0
Gender	0
Churn	0
Outage_sec_perweek	0
Email	0
Contacts	0
Yearly_equip_failure	0
Techie	0
Contract	0
Port_modem	0
Tablet	0
InternetService	0
Phone	0
Multiple	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
PaperlessBilling	0
PaymentMethod	0
Tenure	0
MonthlyCharge	0
Bandwidth_GB_Year	0
Item1	0
Item2	0
Item3	0
Item4	0
Item5	0
Item6	0

```
Item7      0
Item8      0
dtype: int64
```

separate continuous from categorical variables.

```
In [3]: # separate continuous variables  
dfcon = df[['Age','Income','Bandwidth_GB_Year','MonthlyCharge','Outage_sec_perweek','Contacts','Yearly_equip_failure',
```

Standardize continuous variables. Write prepared data to file.

```
In [4]: #standardize data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
data = df
df = dfcon
# scale the data frame
df = scaler.fit_transform(df)
#write the prepared data to .csv file
pd.DataFrame(df).to_csv('prepared-data.csv', index=False)
```

```
/usr/lib/python3/dist-packages/scipy/__init__.py:146: UserWarning: A NumPy version >=1.17.3 and <1.25.0 is required for
this version of SciPy (detected version 1.26.4
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

D

1)

I determined that 2 is the optimal number of clusters using the elbow method. This means that adding more than 2 clusters does not significantly decrease the inertia or within cluster sum of squares variance within the clusters.

2)

Code to plot the inertia of the clusters using the elbow method.

```
In [5]: import matplotlib.pyplot as plt
        from sklearn.cluster import KMeans

        # Define a range of cluster numbers to test
        k values = range(1, 11) # Test cluster numbers from 1 to 10
```

```


inertiaArray = []

# Calculate inertia for each `k` value
for k in k_values:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(df)
    inertiaArray.append(kmeans.inertia_)

# Plot inertia to find the "elbow"
plt.plot(k_values, inertiaArray, 'bo-') # 'bo-' indicates blue circles with lines
plt.xlabel("Number of Clusters (k)")
plt.ylabel("Inertia (Within-Cluster Sum of Squares)")
plt.title("Elbow Method for Determining Optimal Clusters")
plt.show()

## create clusters
kmeans = KMeans(n_clusters=2)
kmeans.fit(df)

```



Out[5]:

KMeans ⓘ ?

KMeans(n_clusters=2)

Inertia

In [6]:

```

inertia = kmeans.inertia_
print(inertia)

```

61704.925171052455

Silhouette Score

```
In [7]: from sklearn.metrics import silhouette_score

silhouette_score = silhouette_score(df, kmeans.labels_)
print(silhouette_score)
```

0.22818537888812102

Davies-Bouldin Index

```
In [8]: from sklearn.metrics import davies_bouldin_score

davies_bouldin_score = davies_bouldin_score(df, kmeans.labels_)
print(davies_bouldin_score)
```

1.7580509340625434

Calinski-Harabasz Index (Variance Ratio Criterion):

```
In [9]: from sklearn.metrics import calinski_harabasz_score

calinski_harabasz_score = calinski_harabasz_score(df, kmeans.labels_)
print(calinski_harabasz_score)
```

2964.336438829804

Print cluster stats

```
In [10]: data['Cluster'] = kmeans.labels_
for cluster_label in data['Cluster'].unique():
    # Subset data for the current cluster
    cluster_data = data[data['Cluster'] == cluster_label]

    # Compute cluster statistics
    cluster_stats = cluster_data.describe()

    # Print cluster statistics
    print(f"Cluster {cluster_label} Statistics:")
    print(cluster_stats)
```

Cluster 1 Statistics:

	CaseOrder	Zip	Lat	Lng	Population \
count	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000
mean	2502.953409	49203.606879	38.834079	-90.829257	9756.015597
std	1447.648689	27695.206694	5.476982	15.263438	14262.714108
min	1.000000	601.000000	17.966120	-171.688150	0.000000
25%	1251.000000	26222.000000	35.458930	-97.156460	722.000000
50%	2502.000000	48836.000000	39.500680	-87.963530	2889.000000
75%	3752.000000	71969.000000	42.120990	-80.003610	13489.000000
max	8572.000000	99927.000000	70.640660	-65.943130	98660.000000

	Children	Age	Income	Outage_sec_perweek \
count	5001.000000	5001.000000	5001.000000	5001.000000
mean	2.095381	52.677465	39737.006721	9.992615
std	2.154507	20.698052	28029.785892	2.977881
min	0.000000	18.000000	348.670000	0.120058
25%	0.000000	35.000000	19287.420000	8.025412
50%	1.000000	52.000000	33377.200000	10.016880
75%	3.000000	71.000000	53517.120000	11.976770
max	10.000000	89.000000	258900.700000	21.207230

	Email	Contacts	Yearly equip_failure	Tenure \
count	5001.000000	5001.000000	5001.000000	5001.000000
mean	12.051590	0.990202	0.392322	9.134829
std	2.988467	0.983109	0.628373	6.041764
min	1.000000	0.000000	0.000000	1.000259
25%	10.000000	0.000000	0.000000	4.332135
50%	12.000000	1.000000	0.000000	7.918063
75%	14.000000	2.000000	1.000000	12.573060
max	23.000000	7.000000	4.000000	37.119120

	MonthlyCharge	Bandwidth_GB_Year	Item1	Item2 \
count	5001.000000	5001.000000	5001.000000	5001.000000
mean	172.701463	1312.214450	3.499500	3.503099
std	42.867453	572.374737	1.033465	1.025495
min	79.978860	155.506715	1.000000	1.000000
25%	139.981600	886.340074	3.000000	3.000000
50%	169.937800	1236.530575	3.000000	4.000000
75%	200.146524	1671.330908	4.000000	4.000000
max	290.160419	3452.422228	7.000000	7.000000

	Item3	Item4	Item5	Item6	Item7 \
count	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000
mean	3.485103	3.503299	3.470706	3.511098	3.511498
std	1.026561	1.031717	1.018819	1.039195	1.021527
min	1.000000	1.000000	1.000000	1.000000	1.000000
25%	3.000000	3.000000	3.000000	3.000000	3.000000
50%	3.000000	4.000000	3.000000	4.000000	4.000000

75%	4.000000	4.000000	4.000000	4.000000	4.000000
max	7.000000	7.000000	7.000000	8.000000	7.000000

	Item8	Cluster
count	5001.000000	5001.0
mean	3.510898	1.0
std	1.033988	0.0
min	1.000000	1.0
25%	3.000000	1.0
50%	4.000000	1.0
75%	4.000000	1.0
max	8.000000	1.0

Cluster 0 Statistics:

	CaseOrder	Zip	Lat	Lng	Population	\
count	4999.000000	4999.000000	4999.000000	4999.000000	4999.000000	
mean	7499.045809	49103.012202	38.681024	-90.735796	9757.109422	
std	1446.149085	27370.827026	5.396952	15.049419	14602.198349	
min	1899.000000	683.000000	18.005430	-170.485200	0.000000	
25%	6249.500000	26358.500000	35.190340	-97.001905	761.000000	
50%	7500.000000	49037.000000	39.292620	-87.908510	2917.000000	
75%	8750.500000	71824.500000	42.077680	-80.194705	12924.500000	
max	10000.000000	99929.000000	70.368530	-65.667850	111850.000000	

	Children	Age	Income	Outage_sec_perweek	\
count	4999.000000	4999.000000	4999.000000	4999.000000	
mean	2.080016	53.479496	39876.874795	10.011085	
std	2.140054	20.694009	28371.726422	2.974425	
min	0.000000	18.000000	643.200000	0.099747	
25%	0.000000	35.000000	19167.905000	8.015255	
50%	1.000000	53.000000	33016.710000	10.020680	
75%	3.000000	71.000000	53007.500000	11.953570	
max	10.000000	89.000000	256998.400000	20.625040	

	Email	Contacts	Yearly_equip_failure	Tenure	\
count	4999.000000	4999.000000	4999.000000	4999.000000	
mean	11.980396	0.998200	0.403681	59.927706	
std	3.062771	0.993877	0.643459	8.479329	
min	1.000000	0.000000	0.000000	31.790270	
25%	10.000000	0.000000	0.000000	54.379310	
50%	12.000000	1.000000	0.000000	61.479870	
75%	14.000000	2.000000	1.000000	66.857995	
max	22.000000	7.000000	6.000000	71.999280	

	MonthlyCharge	Bandwidth_GB_Year	Item1	Item2	\
count	4999.000000	4999.000000	4999.000000	4999.000000	
mean	172.548138	5473.300867	3.482096	3.507101	
std	43.022784	751.908006	1.042144	1.043808	
min	79.978860	3170.023123	1.000000	1.000000	

25%	139.967800	4967.359137	3.000000	3.000000
50%	167.456400	5586.428510	3.000000	4.000000
75%	202.443300	6036.215032	4.000000	4.000000
max	290.160400	7158.981530	7.000000	7.000000

	Item3	Item4	Item5	Item6	Item7 \
count	4999.000000	4999.000000	4999.000000	4999.000000	4999.000000
mean	3.488898	3.491698	3.515103	3.483497	3.507502
std	1.029491	1.019950	1.030411	1.027862	1.035530
min	1.000000	1.000000	1.000000	1.000000	1.000000
25%	3.000000	3.000000	3.000000	3.000000	3.000000
50%	3.000000	3.000000	4.000000	3.000000	4.000000
75%	4.000000	4.000000	4.000000	4.000000	4.000000
max	8.000000	7.000000	7.000000	7.000000	7.000000

	Item8	Cluster
count	4999.000000	4999.0
mean	3.480296	0.0
std	1.023123	0.0
min	1.000000	0.0
25%	3.000000	0.0
50%	3.000000	0.0
75%	4.000000	0.0
max	7.000000	0.0

In []:

E

1)

The quality of the clusters that were created are being evaluated by several different metrics. The clusters have an inertia score of 61704.92517105245. A lower inertia indicates a tighter cluster. The clusters have a silhouette score of 0.22818537888812102. Positive silhouette scores indicate that the data points are closer to their own cluster center than to other cluster centers. A higher silhouette score indicates better clustering. The clusters have a davies-bouldin score of 1.7580509340625434. A lower davies-bouldin indicates well separated and compact clusters. The clusters have a Calinski-Harabasz Index score of 2964.336438829804. A higher Calinski-Harabasz score indicates better defined clusters. Based on the values of the metrics stated above I think the quality of the cluster analysis is very good. There is nothing indicating poor clustering such as a negative silhouette score.

2)

The results of the cluster analysis show that the customers in the churn data set can be grouped into 2 clusters. The number of clusters was dictated by using the elbow method against inertia values of different numbers of clusters. The analysis shows that the quality of the clusters is good so there are statistically significant differences in the characteristics of the customers assigned to each cluster. The cluster sizes are almost equal.

Some observations from each cluster are that the mean population of cluster 1 is slightly greater than cluster 0. The mean age of customers in cluster 1 is slightly greater than cluster 0. The mean income of customers in cluster 1 is slightly greater than that of customers in cluster 0. The mean outage_sec_perweek is slightly greater in cluster 1 compared to cluster 0. The mean monthly charge is slightly greater in cluster 0 compared to cluster 1.

Some implications of this cluster analysis are that the company can use this information to better serve customers in each cluster since the customers assigned to each cluster have different characteristics. For instance the mean outage_sec_perweek is greater in cluster 1. This implies that customers in that cluster should be targeted for services like online backup and should also be targeted for better service reliability to reduce customer churn. Another implication is that customers in cluster 0 have a higher monthly charge and therefore are good candidates for sales promotions that are designed by the marketing team to reduce customer churn.

3)

One limitation of my k-means cluster analysis is that the analysis only uses numeric continuous variables. I think that it would more informative if the clustering could use categorical variables such as 'Churn'

and 'Streaming_TV'. Because of this there is a lot of customer data not being used to form the clusters.

4)

One course of action based on the results of this analysis would be that we should spend more time and resources increasing reliability of services for customers assigned to cluster 1 because of the higher mean outage_sec_perweek. Sales and promotional campaigns can be tailored to reduce the monthly charge for customers assigned to cluster 0 because of the higher mean monthly charge. Advertisements can be created for a slightly older demographic and presented to the customers assigned to cluster 1 because of the higher mean age in that cluster.

In []: